

RISC-V ISA Profile of the XMSS C Implementation

XMSS-Jasmin Project

February 2026

1 Context

XMSS (RFC 8391) is a stateful hash-based signature scheme standardised for post-quantum use. We are building a formally verifiable implementation in Jasmin, targeting x86-64 first (where the Jasmin compiler is mature) and RISC-V second (where the backend is under active development).

Before writing RISC-V Jasmin code—or contributing to the Jasmin RISC-V backend—we need to know which ISA extensions XMSS actually exercises. This report answers that question by disassembling the compiled C implementation and classifying every instruction.

The key question is: *does XMSS require anything beyond RV64I + M?* If the algorithm logic is pure I+M, then ISA extension work (Zbb rotates for SHA-2, etc.) is confined to the hash layer and does not affect the algorithm-level Jasmin code.

2 Methodology

2.1 Analysis target

We analyse `libxmss.a`—the static library containing only XMSS algorithm code: parameter derivation, hash wrappers, WOTS+, L-tree, treehash, BDS state management, XMSS signing/verification, and XMSS-MT. This comprises 13 object files.

Crucially, the library excludes `printf`, `malloc`, stack guards, and other libc/test-harness code. Analysing test binaries (the approach of a previous iteration) would have included hundreds of libc instructions that have nothing to do with XMSS.

2.2 Toolchain

Compiler	<code>riscv64-linux-gnu-gcc</code> 13.3.0 (Ubuntu)
Flags	<code>-march=rv64gc -mabi=lp64d -O3</code>
Disassembler	<code>riscv64-linux-gnu-objdump</code> (Binutils 2.42)

The `rv64gc` march is the standard general-purpose profile: RV64I + M + A + F + D + Zicsr + Zifencei + C. It does *not* include Zba, Zbb, or other bitmanip extensions.

2.3 Classification

Each disassembled instruction is classified on two orthogonal axes:

1. **Semantic extension** (what the instruction does): I, M, A, F, D, Zba, Zbb, etc. Determined by looking up the mnemonic in a table generated from the `riscv-opcodes` database (the same database used by the RISC-V toolchain).
2. **Encoding width** (how it is encoded): 16-bit compressed (C extension) or 32-bit standard. Determined from the raw instruction byte count in the `objdump` output.

The semantic extension is what matters for Jasmin. C encoding is a secondary observation: the assembler handles it automatically and it does not affect which instructions Jasmin source code must express.

A subtlety: GNU `objdump` renders compressed instructions using their uncompressed aliases (`sd` not `c.sd`). A naïve classifier that looks for `c.` prefixes would report C = 0%. We detect C encoding from byte width instead.

3 Findings

3.1 Extension summary

Table 1 shows the complete picture: XMSS uses only the I and M extensions.

Table 1: Semantic extension summary across all 13 object files.

Extension	Instructions	Unique mnemonics	Notes
I	9505	46	Base integer: loads, stores, branches, arithmetic, shifts
M	57	3	<code>mulw</code> (37), <code>mul</code> (17), <code>divuw</code> (3)
A, F, D	0	0	Not present despite being enabled by <code>rv64gc</code>
Zba, Zbb, Zbs	0	0	Not present (not in <code>rv64gc</code> ; see §4)

99.4% of instructions are base integer (I). The 0.6% that are M come entirely from compiler-generated address arithmetic: `mulw` for array index calculations, `mul` for 64-bit offset computation, and `divuw` for parameter derivation in `params.c`.

3.2 Per-module breakdown

Table 2 shows instruction counts per object file, grouped into algorithm modules and hash modules.

Two observations:

- The hash layer (SHA-2, SHAKE, hash dispatch) uses *zero* M instructions. It is pure RV64I.
- M instructions appear only in the algorithm layer, and only for index/offset arithmetic the compiler generates from C expressions like `idx * n` or `h / d`.

3.3 Compressed encoding

48% of instructions use 16-bit C encoding. This varies from 33% (`sha2_local.c.o`, which has many 32-bit shift-immediate instructions for SHA-2 rotations) to 69% (`ltree.c.o`, which is mostly register moves and branches).

Table 2: Per-object-file instruction counts. M column shows M-extension instruction count; all remaining instructions are I.

Object file	Total	M insns	C-encoded (%)
<i>Hash layer</i>			
<code>sha2_local.c.o</code>	1953	0	33%
<code>shake_local.c.o</code>	1337	0	44%
<code>xmss_hash.c.o</code>	1082	0	51%
<i>Algorithm layer</i>			
<code>bds.c.o</code>	1130	5	55%
<code>xmss_mt.c.o</code>	1145	18	58%
<code>wots.c.o</code>	760	4	47%
<code>xmss.c.o</code>	637	3	60%
<code>treehash.c.o</code>	484	2	50%
<code>bds_serialize.c.o</code>	473	14	53%
<code>params.c.o</code>	286	6	60%
<code>ltree.c.o</code>	129	5	69%
<code>address.c.o</code>	99	0	38%
<code>utils.c.o</code>	47	0	68%
Total	9562	57	48%

C encoding is handled automatically by the assembler and is invisible to Jasmin source code. It reduces code size but does not affect correctness or the set of required ISA extensions.

4 The Zbb question

The most interesting finding is what is *absent*. The hash modules (`sha2_local.c.o`, `shake_local.c.o`) implement SHA-256, SHA-512, and Keccak using only RV64I operations. SHA-2 in particular requires 32-bit rotations, which the compiler synthesises as shift-shift-or sequences:

```
srliw a5, a5, 17    # high part
slliw a4, a4, 15    # low part
or     a5, a5, a4    # combine
```

The Zbb extension provides a single `rорw` instruction that replaces this 3-instruction sequence. Similarly, `rev8` (Zbb) would replace multi-instruction byte-swap sequences in SHA-2 endianness conversion.

Since `rv64gc` does not include Zbb, the compiler cannot emit these instructions. But a Jasmin implementation targeting `rv64gc_zbb` could use them explicitly in the hash layer, gaining performance without affecting the algorithm layer at all.

This reinforces the algorithm/hash boundary in the C implementation's architecture: the hash layer is the only place where ISA-specific optimisation is relevant. All algorithm-layer Jasmin code can be written in pure RV64I (with compiler-inserted M for index arithmetic).

5 Recommended Jasmin targets

rv64im (minimum)

Base integer + multiply/divide. Sufficient for all XMSS algorithm logic. The hash layer would use software rotations.

rv64gc (standard)

The standard Linux general-purpose profile. Adds A, F, D, Zicsr, Zifencei, and C—none of which XMSS uses, but targeting it ensures binary compatibility with standard RISC-V Linux distributions.

rv64gc_zbb (optimised)

Adds Zbb to the standard profile. The hash layer can exploit `rorw/ror` for SHA-2 rotations and `rev8` for endianness conversion. Algorithm-layer code is unchanged.

6 Next steps

1. The Jasmin implementation starts with x86-64, where the compiler backend is mature and the hash layer can use native rotation instructions (`ROR`).
2. Algorithm-layer Jasmin code (WOTS+, BDS, treehash, XMSS sign/verify) should use only portable basic operations—no architecture-specific intrinsics. The ISA analysis confirms this is sufficient.
3. When the Jasmin RISC-V backend matures, the algorithm layer ports directly. The hash layer is the only component requiring architecture-specific work (Zbb rotations).